CS1632: TESTING THEORY AND TERMINOLOGY

Wonsun Ahn

Expected vs. Observed Behavior

- **Expected behavior**: What "should" happen
- **Observed behavior**: What "does" happen

- **Testing**: checking expected == observed behavior
- **Defect**: a flaw that causes expected != observed behavior

- Expected behavior is enforced using **requirements**

Example

- Suppose we are testing a function `sqrt`:
`// Requirement: return square root of num`
`float sqrt(int num) { … }`

- For `ret = sqrt(9)`, **expected behavior**: `ret == 3`
  → Any **observed behavior** where `ret != 3` is a **defect**

- For `ret = sqrt(-9)`, what is the **expected behavior**?
  → Depends, but **requirements** should specify some behavior (e.g., return 0 for all negative numbers)
  → Any **observed behavior** other than the above is a **defect**

THE IMPOSSIBILITY OF EXHAUSTIVE TESTING

- Let's say we want to ensure that `sqrt` is defect-free for all arguments (both positive and negative)

- Assume arg is a Java `int` (signed 32-bit integer)

- How many values do we have to test?

4,294,967,296

What if there are two arguments?

- Suppose we are testing a function `add`:
  ```
  // return the sum of x and y
  int add(int x, int y) { … }
  ```

- How many tests do we have to perform?
  (Hint: all combinations of x and y)

4,294,967,296 ^ 2

What if the argument is an array?

- Suppose we are testing a function `add`:
  ```
  // return sum of elements in A
  int add(int[] A) { … }
  ```

- How many tests do we have to perform? (Note: array `A` can be arbitrarily long)

4,294,967,296 < Infinity

Would testing all the combinations of arguments guarantee that there are no problems?

LOL NOPE

- Issues not covered by exhaustive input testing
  - Compiler issues
  - Systems-level issues (e.g. OS/device-dependent defect)
  - Parallel programming issues (e.g. data races)

- The same input must be tested multiple times
  - On different compilers, OSes, devices, …
  - (Potentially) many times on same compiler / OS / device

Compiler Issues

- The compiled binary, not your source code, runs on the computer
- What if compiler has a bug? (Rare)
- What if compiler *exposes* a bug in your program? (More frequent)

```
int add_up_to (int count) {
    int sum, i;      /* Is sum == 0? Not necessarily! */
    for(i = 0; i <= count; i++) sum = sum + i;
    return sum;
}
```
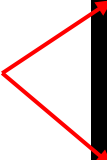
- Behavior is undefined according to C language specifications
- Compiler can generate code with arbitrary behavior

☞ Need to exhaustively verify with all compilers and compiler options!

Parallel programming issues

```java
class Main implements Runnable {
    public static int count = 0;
    public void run() {
        for(int i=0; i < 1000000; i++) { count++; }
        System.out.println("count = " + count);
    }
    public static void main(String[] args) {
        Main m = new Main();
        Thread t1 = new Thread(m);
        Thread t2 = new Thread(m);
        t1.start();
        t2.start();
    }
}
```

Why?

```
$ javac Main.java
$ java Main
count = 1868180
count = 1868180
$ java Main
count = 1033139
count = 1033139
```

Parallel programming issues

- Why does this happen?
  - Threads `t1` and `t2` execute concurrently
  - Two threads try to increment `count` at the same time
  - Often, they step on each other's toes (a *data race*)
- If there is a data race, result is undefined
  - Java language specifications say so!
  - Every time you run it, you may get a different result
  - Result depends on relative speed of threads `t1` and `t2`
- ☞ Running 1000+ times may not cover all behavior

For the purposes of this Chapter…

- Let's ignore these issues for now
  - Combinatorial testing issues
  - Compiler issues
  - Systems-level issues
  - Parallel programming issues

- Exhaustive input value testing is hard enough
  - a.k.a. "test explosion problem"
  - This is what we will focus on in this chapter

- We will address the other issues later 😉

Equivalence Classes

Achieving Test Coverage Efficiently

Defining Test Coverage

- Goal of testing: achieve good test coverage
  - **Test coverage**: measure of how well code has been tested
  - Ideally, *test_coverage = defects_found / total_defects*

- But is there a way to measure *total_defects*?
  - If we knew, we wouldn't need to do any testing!
  - Impossible to measure true test coverage

- Is there a good proxy that estimates true test coverage?
  - **Statement coverage** = *statements_tested / total_statements*
  - Rationale: if a high percentage of statements are tested
    ☞ likely that a high percentage of defects are found
  - Other proxies out there: method coverage, path coverage, …

Improving Test Coverage

- QA engineers have a limited testing time budget
    - Must choose tests maximizing test coverage

- Which tests are likely to maximize coverage?
    - Tests that exercise all **required** program behaviors
    - If tests focus only on a subset of program behaviors $\rightarrow$ cannot find defects in behaviors that are not covered
    - This is the idea behind *equivalence class partitioning*

Equivalence Class Partitioning

- Partition the input values into "equivalence classes"
  - Equivalence class = group of values with same behavior

- E.g. equivalence classes for our `sqrt` method: {*nonnegative_numbers, negative_numbers*}

- Behavior for each equivalence class:
  - *nonnegative_numbers*: returns square root of number
  - *negative_numbers*: returns 0

Equivalence Classes should be *Strictly* Partitioned

- **Strictly**: each value belongs to one and only one class

- If an input value belongs to multiple classes
  - Means requirements specify two different behaviors for the input
  - Either requirements are inconsistent, or you misunderstood them

- If an input value belongs to no class
  - Means requirements do not specify a behavior for the input
  - Either requirements are incomplete, or you misunderstood them

Values can be Strings

- For a spell checker, input values are strings

- Equivalence classes:
  {*strings_in_dictionary*, *strings_not_in_dictionary*}

- Behavior for each equivalence class:
  - *strings_in_dictionary*: do nothing
  - *strings_not_in_dictionary*: red underline string

Values can also be Objects

- Input values can be tuna cans

- Equivalence classes:
  {*not_expired, expired_and_not_smelly, expired_and_smelly*}

- Behavior for each equivalence class:
  - *not_expired*: eat
  - *expired_and_not_smelly*: feed it to your cat (kidding)
  - *expired_and_smelly*: discard

22

Testing Each Equivalence Class

- Pick at least one value from each equivalence class
  - Ensures you cover all behavior expected of program
  - Gets you good coverage without exhaustive testing!

- How many values should I pick? And what values?
  - There is no exact science.
  - But there are some good empirical guidelines!

Defects are Prevalent at Boundaries

- Empirical truth:
  - Defects are more prevalent at boundaries of equivalence classes than in the middle.

- Why?
  - Due to prevalence of **off-by-one** errors

Case Study: Choosing Input Values

- Suppose the requirements are:

1. Age shall be given as a commandline argument

2. If given age is equal to or less than **34 years**, the system shall print "cannot be US president".

3. If given age is equal to or greater than **35 years**, the system shall print "can be US president".

Equivalence class partitioning

CANNOT_BE_PRESIDENT =
[...19,20,21,22,23,24,25,26,27,28,29,30,31,32,33,34]

CAN_BE_PRESIDENT =
[35,36,37,38,39,40,41,42,43,44,45,46,47,48,49,50...]

Always Test Boundary Values

CANNOT_BE_PRESIDENT =
[…19,20,21,22,23,24,25,26,27,28,29,30,31,32,33,34]

CAN_BE_PRESIDENT =
[35,36,37,38,39,40,41,42,43,44,45,46,47,48,49,50…]

- Always test boundary values (shown in red).
- Will catch off-by-one errors at 34 or 35. E.g.,

```
if (age > 35) {ret = "can be US president";}
```

Also Test a few Interior Values

CANNOT_BE_PRESIDENT =
[...19,20,21,22,23,24,25,26,27,28,29,30,31,32,33,34]

CAN_BE_PRESIDENT =
[35,36,37,38,39,40,41,42,43,44,45,46,47,48,49,50...]

- Testing interior values (in green) is also important.
- Who knows? Code behavior may change in the middle.

Are we done?

CANNOT_BE_PRESIDENT =
[...19,20,21,22,23,24,25,26,27,28,29,30,31,32,33,34]

CAN_BE_PRESIDENT =
[35,36,37,38,39,40,41,42,43,44,45,46,47,48,49,50...]

- Test input values so far: {26, 30, 34, 35, 39, 42}

Implicit (hidden) boundary values

- **Explicit boundaries**: specified in requirements
  - Boundaries that are defined by equivalence classes
- **Implicit boundaries**: not in the requirements
  - "Naturally" occurring in language, hardware, domain
  - Language boundaries: MAXINT, MININT
  - Hardware boundaries: memory space, hard drive space
  - Domain boundaries: weight >= 0, 0 <= score <= 100, etc.

- We need to check implicit boundaries as well. Why?

Implicit boundaries should not change behavior

- Why do we check **explicit boundaries**?
  - To verify behavior **changes** when boundary is crossed.
- Why do we check **implicit boundaries**?
  - To verify behavior **does not change** on the boundary.

- Checking correct handling of implicit boundaries:
  - Crossing MAXINT boundary for input value
    $\rightarrow$ Handle larger numbers gracefully with no int overflow
  - Crossing memory space limit due to large input
    $\rightarrow$ Handle gracefully (possibly by moving data to disk)

Add implicit boundary values

CANNOT_BE_PRESIDENT =
[MININT-1,MININT,…,-1,0,1,…,26,27,28,29,30,31,32,33,34]

CAN_BE_PRESIDENT =
[35,36,37,38,39,40,41,42,43,44,…,MAXINT,MAXINT+1]

- language boundaries: MININT <= age <= MAXINT
- domain boundary: age >= 0 (age is typically non-negative)
- Inputs: {MININT-1, MININT, -1, 0, 26, 30, 34, 35, 39, 42, MAXINT, MAXINT+1}

Finding Defects using our Inputs

- Now, let's feed the inputs to our code:
  {MININT-1, MININT, -1, 0, 26, 30, 34, 35, 39, 42, MAXINT, MAXINT+1}

```
public static void main(String[] args) {
  int age = Integer.parseInt(args[0]);
  if (age > 35) {   // The off-by-one error.
    System.out.println("can be US president");
  } else {
    System.out.println("cannot be US president");
  }
}
```

- Which inputs would find defects?

Finding Defects using our Inputs

- Defect 1: off-by-one-error found with input 35:
  - Expected behavior: prints "can be US president"
  - Observed behavior: prints "cannot be US president"

- Defect 2: int overflow error found with input MAXINT+1:
  - Expected behavior: prints "can be US president"
  - Observed behavior: throws java.lang.NumberFormatException

- Defect 3: int overflow error found with input MININT-1:
  - Expected behavior: prints "cannot be US president"
  - Observed behavior: throws java.lang.NumberFormatException

Base, edge, and corner cases

- **Base case**: A typical use case
  - Interior value of equivalence class for normal operation

- **Edge case**: A use case at the limit of allowed use
  - Boundary value of equivalence class for normal operation

- **Corner case** (or **pathological case**):
  - Multiple edge cases happening simultaneously
  - Or value far outside of normal operating parameters

Example: Base, edge, corner cases

- Suppose a cat scale has these operating envelopes:
  - Weight between 0 – 100 lbs
  - Temperature between 0 – 120 F
- Base cases: (10 lbs, 60 F), (20 lbs, 70 F), …
- Edge cases: (**100 lbs**, 70 F), (10 lbs, **0 F**), …
- Corner cases: (**10000 lbs**, 70 F), (**100 lbs**, **120 F**), …
- Why test 10000 lbs?
  - Even if scale isn't expected to operate correctly for 10000 lbs, user still cares what happens (e.g. shouldn't break the scale)

Categories of Testing:
Black / White / Gray
Dynamic / Static

Black-, white, and gray-box testing

- **Black-box testing**:
  - Testing with no knowledge of interior structure source code
  - Tests are performed from the user's perspective
  - Can be performed by lay people who don't know how to program

- **White-box testing**:
  - Testing with explicit knowledge of the interior structure and codebase
  - Tests are performed from the developer's perspective
  - Test inputs are crafted to exercise specific lines of code

- **Gray-box testing**:
  - Testing with some knowledge of the interior structure and codebase
  - Knowledge comes from partial code inspection or a design document
  - Performed from the user's perspective, but informed by knowledge

Black-box testing examples

- Tests are performed using only UI

- Examples:
  - Testing a website using a web browser
  - Testing a game by actually playing it
  - Testing a script against an API endpoint
  - Any type of beta test
  - Penetration testing on a website

White-box testing examples

- Tests are performed by both…
  - Using UI to exercise specific program paths
  - Explicitly calling methods from a testing script
- Examples
  - Choosing inputs to exercise specific parts of an algorithm
  - Choosing inputs causing exceptions and checking handling
  - Testing that a method call returns the correct result
  - Testing that instantiating a class creates a valid object
- Unlike black-box, can measure **statement coverage**

Static vs dynamic testing

- We talked a great deal about choosing good inputs
  - But is this all there is to testing?

- Dynamic testing = code is executed
  - Relies on good inputs for good coverage

- Static testing = code is not executed
  - There are no inputs since code is not executed
  - Relies on analyzing the code to find defects

Dynamic testing

- Code is executed under various test scenarios
  - Varying input values, compilers, OSes, etc.
  - Observed results are compared with expected results
  - Hard to achieve 100% test coverage

- Examples:
  - Manual testing
  - Unit testing
  - System testing
  - Performance testing

Static testing

- Code is analyzed by a person or testing tool
  - While checking whether correctness rules are followed
  - 100% test coverage achieved for all code analyzed
  - Even when check passes, defects can still occur at runtime

- Examples:
  - Code reviews by a person
  - Code analysis using a tool
    - Compilers
    - Linters
    - Bug pattern finders
    - Code coverage analysis
    - Model checkers

Now Please Read Textbook Chapters 2-4